

AV-HuBERT: Self-Supervised Learning of Audio-Visual Speech Representation

Wei-Ning Hsu

Meta AI

Joint work with Bowen Shi, Kushal Lakhotia, Abdelrahman Mohamed

Feb 28, 2022 @ AAI workshop on Self-Supervised Learning for Audio and Speech Processing

Motivation

Automatic speech recognition (ASR) is widely used, BUT

- Performance degrades a lot when noisy
- Especially when the noise is speech
- On LRS3, from 4.7% WER (clean) to 32.1% WER (0 dB)

We need to make ASR more robust, how?

- **Audio-visual** speech recognition (AVSR)
 - Use complementary visual information (lip)
 - Invariant to noise
 - In the case of $-\infty$ SNR / no audio \rightarrow lip-reading (VSR)
- VSR is also useful for people with speech impairment

Motivation

Automatic speech recognition (ASR) is widely used, BUT

- Performance degrades a lot when noisy
- Especially when the noise is speech
- On LRS3, from 4.7% WER (clean) to 32.1% WER (0 dB)


We need to make ASR more robust, how?

- **Audio-visual** speech recognition (AVSR)
 - Use complementary visual information (lip)
 - Invariant to noise
 - In the case of $-\infty$ SNR / no audio \rightarrow lip-reading (VSR)
- VSR is also useful for people with speech impairment



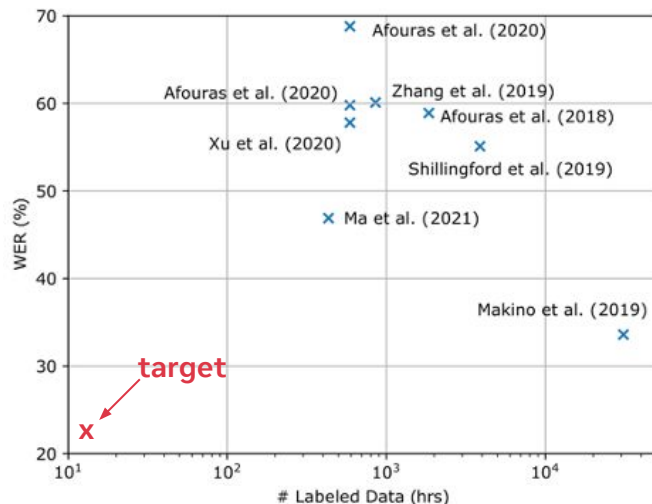
Motivation (cont.)

- What are the barriers?
 - Lack of labeled data
 - Largest public datasets: LRS2* (224), LRS3 (433)
 - Far less than ASR data: Librispeech (1K), GigaSpeech (10K)
 - VSR and AVSR are also data hungry
 - Trained on LRS3: 46.9% WER
 - Trained on YT31k: 33.6% WER
 - Complicated pipeline
 - First pre-trained on isolated words (LRW)
 - Then do curriculum training (short to long)




Motivation (cont.)

- How do we approach that?
 - Self-supervised learning
 - First, pre-train on unlabeled data
 - Then, fine-tune on (limited) labeled data
 - Are unlabeled data available for VSR/AVSR?
 - Yes! news, movies, social media, meeting, ...
 - What self-supervised learning method to use?
 - Adapt from speech SSL




Self-Supervised Learning for Visual Speech Recognition

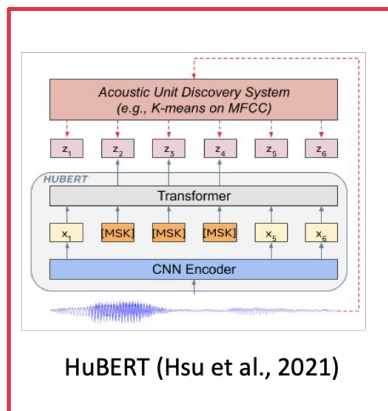
- What is lip-reading (VSR)?
 - Input: a sequence of image frames
 - output: a sequence of characters/word-pieces
 - Supervised learning: trained on (video, text) pairs
- What speech SSL methods to adapt from?



CPC (Oord et al., 2018)

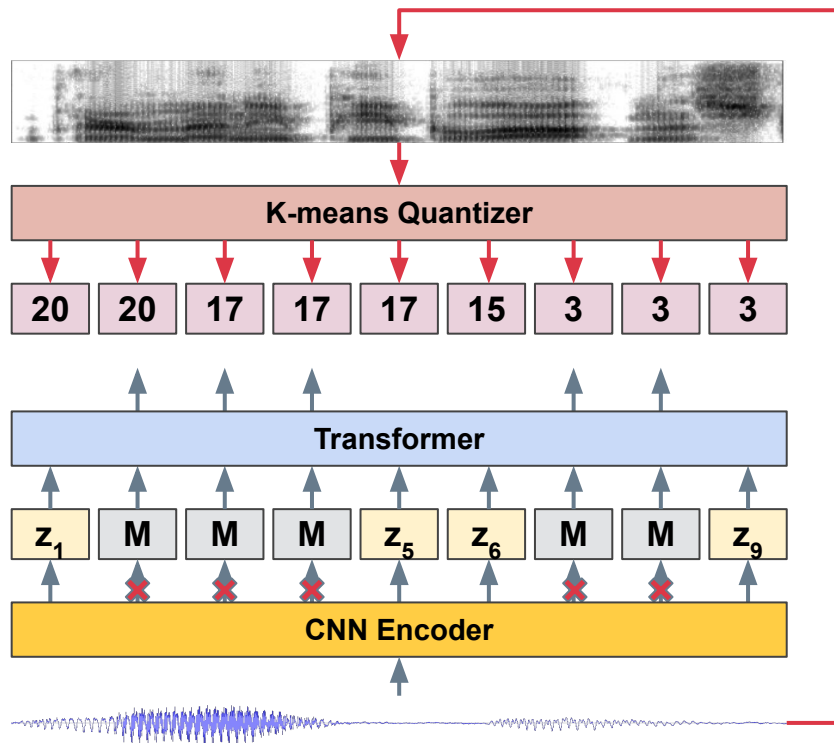


wav2vec 2.0 (Baevski et al., 2021)



HuBERT (Hsu et al., 2021)

Preliminary: HuBERT for audio (A-HuBERT)

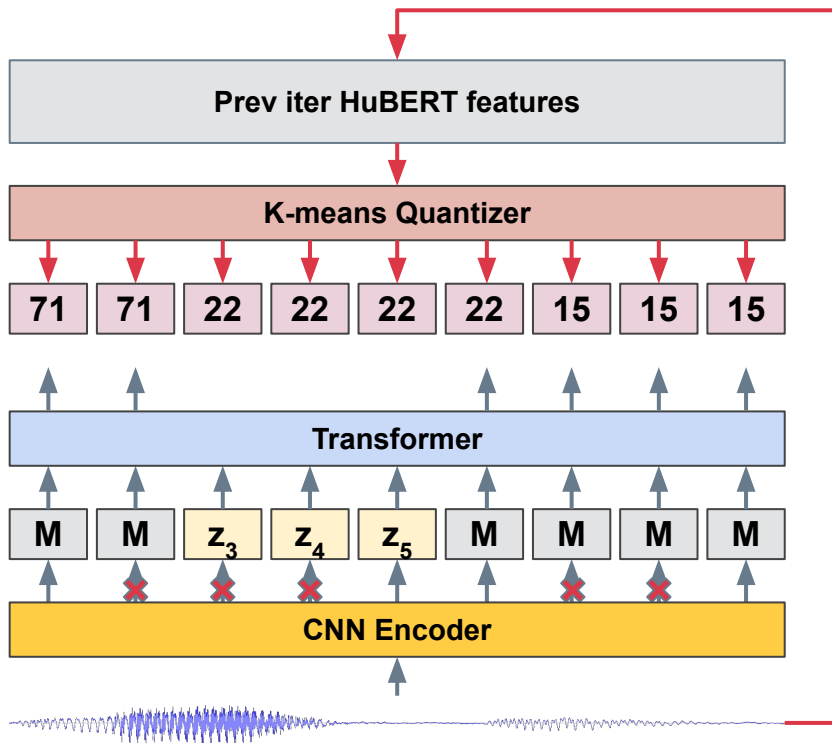


HuBERT is a SOTA SSL framework for speech that performs well on many tasks

How does it work?

- Given an audio stream
 - K-means on MFCC
 - Masked prediction

Preliminary: HuBERT for audio (A-HuBERT)

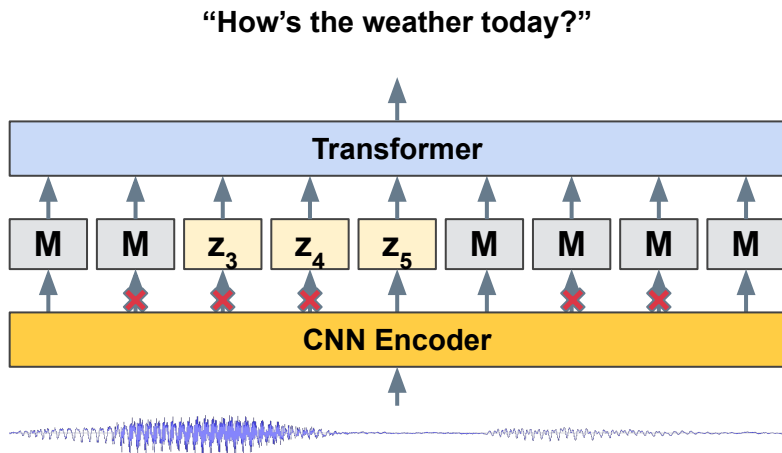


HuBERT is a SOTA SSL framework for speech that performs well on many tasks

How does it work?

- Given an audio stream
 - K-means on MFCC
 - Masked prediction
 - Iterative refinement

Preliminary: HuBERT for audio (A-HuBERT)




HuBERT is a SOTA SSL framework for speech that performs well on many tasks

How does it work?

- Given an audio stream
 1. K-means on MFCC
 2. Masked prediction
 3. Iterative refinement
 4. Remove cluster prediction head and fine-tune with labeled data


Preliminary: HuBERT for audio (A-HuBERT)

- Shown effective for inference and generative tasks
 - ASR



Preliminary: HuBERT for audio (A-HuBERT)


- Shown effective for inference and generative tasks
 - ASR
 - SUPERB (content, speaker, semantic, emotion)




	PR	KS	IC	SID	ER	ASR (WER)		QbE	SF		ASV	SD
	PER ↓	Acc ↑	Acc ↑	Acc ↑	Acc ↑	w/o ↓	w/ LM ↓	MTWV ↑	F1 ↑	CER ↓	EER ↓	DER ↓
FBANK	82.01	8.63	9.10	8.5E-4	35.39	23.18	15.21	0.0058	69.64	52.94	9.56	10.05
PASE+ [16]	58.87	82.54	29.82	37.99	57.86	25.11	16.62	0.0072	62.14	60.17	11.61	8.68
APC [7]	41.98	91.01	74.69	60.42	59.33	21.28	14.74	0.0310	70.46	50.89	8.56	10.53
VQ-APC [32]	41.08	91.11	74.48	60.15	59.66	21.20	15.21	0.0251	68.53	52.91	8.72	10.45
NPC [33]	43.81	88.96	69.44	55.92	59.08	20.20	13.91	0.0246	72.79	48.44	9.4	9.34
Mockingjay [8]	70.19	83.67	34.33	32.29	50.28	22.82	15.48	6.6E-04	61.59	58.89	11.66	10.54
TERA [9]	49.17	89.48	58.42	57.57	56.27	18.17	12.16	0.0013	67.50	54.17	15.89	9.96
modified CPC [34]	42.54	91.88	64.09	39.63	60.96	20.18	13.53	0.0326	71.19	49.91	12.86	10.38
wav2vec [12]	31.58	95.59	84.92	56.56	59.79	15.86	11.00	0.0485	76.37	43.71	7.99	9.9
vq-wav2vec [13]	33.48	93.38	85.68	38.80	58.24	17.71	12.80	0.0410	77.68	41.54	10.38	9.93
wav2vec 2.0 Base [14]	5.74	96.23	92.35	75.18	63.43	6.43	4.79	0.0233	88.30	24.77	6.02	6.08
wav2vec 2.0 Large [14]	4.75	96.66	95.28	86.14	65.64	3.75	3.10	0.0489	87.11	27.31	5.65	5.62
HuBERT Base [35]	5.41	96.30	98.34	81.42	64.92	6.42	4.79	0.0736	88.53	25.20	5.11	5.88
HuBERT Large [35]	3.53	95.29	98.76	90.33	67.62	3.62	2.94	0.0353	89.81	21.76	5.98	5.75

Preliminary: HuBERT for audio (A-HuBERT)

- Shown effective for inference and generative tasks
 - ASR
 - SUPERB (content, speaker, semantic, emotion)
 - Speech-to-speech translation



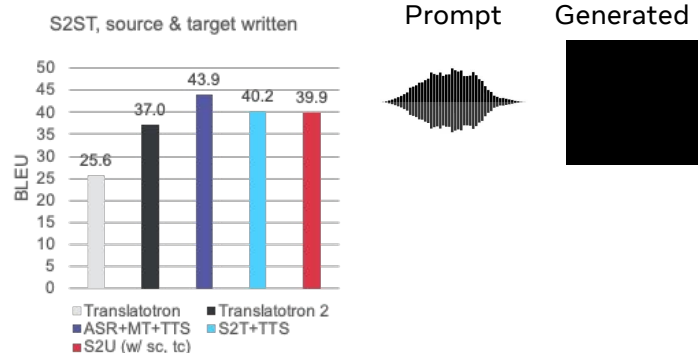
S2ST, source & target written



	PR	KS	IC	SID	ER	ASR (WER)		QbE	SF		ASV	SD
	PER ↓	Acc ↑	Acc ↑	Acc ↑	Acc ↑	w/o ↓	w/ LM ↓	MTWV ↑	F1 ↑	CER ↓	EER ↓	DER ↓
FBANK	82.01	8.63	9.10	8.5E-4	35.39	23.18	15.21	0.0058	69.64	52.94	9.56	10.05
PASE+ [16]	58.87	82.54	29.82	37.99	57.86	25.11	16.62	0.0072	62.14	60.17	11.61	8.68
APC [7]	41.98	91.01	74.69	60.42	59.33	21.28	14.74	0.0310	70.46	50.89	8.56	10.53
VQ-APC [32]	41.08	91.11	74.48	60.15	59.66	21.20	15.21	0.0251	68.53	52.91	8.72	10.45
NPC [33]	43.81	88.96	69.44	55.92	59.08	20.20	13.91	0.0246	72.79	48.44	9.4	9.34
Mockingjay [8]	70.19	83.67	34.33	32.29	50.28	22.82	15.48	6.6E-04	61.59	58.89	11.66	10.54
TERA [9]	49.17	89.48	58.42	57.57	56.27	18.17	12.16	0.0013	67.50	54.17	15.89	9.96
modified CPC [34]	42.54	91.88	64.09	39.63	60.96	20.18	13.53	0.0326	71.19	49.91	12.86	10.38
wav2vec [12]	31.58	95.59	84.92	56.56	59.79	15.86	11.00	0.0485	76.37	43.71	7.99	9.9
vq-wav2vec [13]	33.48	93.38	85.68	38.80	58.24	17.71	12.80	0.0410	77.68	41.54	10.38	9.93
wav2vec 2.0 Base [14]	5.74	96.23	92.35	75.18	63.43	6.43	4.79	0.0233	88.30	24.77	6.02	6.08
wav2vec 2.0 Large [14]	4.75	96.66	95.28	86.14	65.64	3.75	3.10	0.0489	87.11	27.31	5.65	5.62
HuBERT Base [35]	5.41	96.30	98.34	81.42	64.92	6.42	4.79	0.0736	88.53	25.20	5.11	5.88
HuBERT Large [35]	3.53	95.29	98.76	90.33	67.62	3.62	2.94	0.0353	89.81	21.76	5.98	5.75

Preliminary: HuBERT for audio (A-HuBERT)


- Shown effective for inference and generative tasks
 - ASR
 - SUPERB (content, speaker, semantic, emotion)
 - Speech-to-speech translation
 - Textless NLP (<https://speechbot.github.io>):
 - text-free audio GPT




	PR	KS	IC	SID	ER	ASR (WER)		QbE	SF		ASV	SD
	PER ↓	Acc ↑	Acc ↑	Acc ↑	Acc ↑	w/o ↓	w/ LM ↓	MTWV ↑	F1 ↑	CER ↓	EER ↓	DER ↓
FBANK	82.01	8.63	9.10	8.5E-4	35.39	23.18	15.21	0.0058	69.64	52.94	9.56	10.05
PASE+ [16]	58.87	82.54	29.82	37.99	57.86	25.11	16.62	0.0072	62.14	60.17	11.61	8.68
APC [7]	41.98	91.01	74.69	60.42	59.33	21.28	14.74	0.0310	70.46	50.89	8.56	10.53
VQ-APC [32]	41.08	91.11	74.48	60.15	59.66	21.20	15.21	0.0251	68.53	52.91	8.72	10.45
NPC [33]	43.81	88.96	69.44	55.92	59.08	20.20	13.91	0.0246	72.79	48.44	9.4	9.34
Mockingjay [8]	70.19	83.67	34.33	32.29	50.28	22.82	15.48	6.6E-04	61.59	58.89	11.66	10.54
TERA [9]	49.17	89.48	58.42	57.57	56.27	18.17	12.16	0.0013	67.50	54.17	15.89	9.96
modified CPC [34]	42.54	91.88	64.09	39.63	60.96	20.18	13.53	0.0326	71.19	49.91	12.86	10.38
wav2vec [12]	31.58	95.59	84.92	56.56	59.79	15.86	11.00	0.0485	76.37	43.71	7.99	9.9
vq-wav2vec [13]	33.48	93.38	85.68	38.80	58.24	17.71	12.80	0.0410	77.68	41.54	10.38	9.93
wav2vec 2.0 Base [14]	5.74	96.23	92.35	75.18	63.43	6.43	4.79	0.0233	88.30	24.77	6.02	6.08
wav2vec 2.0 Large [14]	4.75	96.66	95.28	86.14	65.64	3.75	3.10	0.0489	87.11	27.31	5.65	5.62
HuBERT Base [35]	5.41	96.30	98.34	81.42	64.92	6.42	4.79	0.0736	88.53	25.20	5.11	5.88
HuBERT Large [35]	3.53	95.29	98.76	90.33	67.62	3.62	2.94	0.0353	89.81	21.76	5.98	5.75

Preliminary: HuBERT for audio (A-HuBERT)

- Shown effective for inference and generative tasks
 - ASR
 - SUPERB (content, speaker, semantic, emotion)
 - Speech-to-speech translation
 - Textless NLP (<https://speechbot.github.io>):
 - text-free audio GPT, emotion conversion



S2ST, source & target written



Prompt

Generated



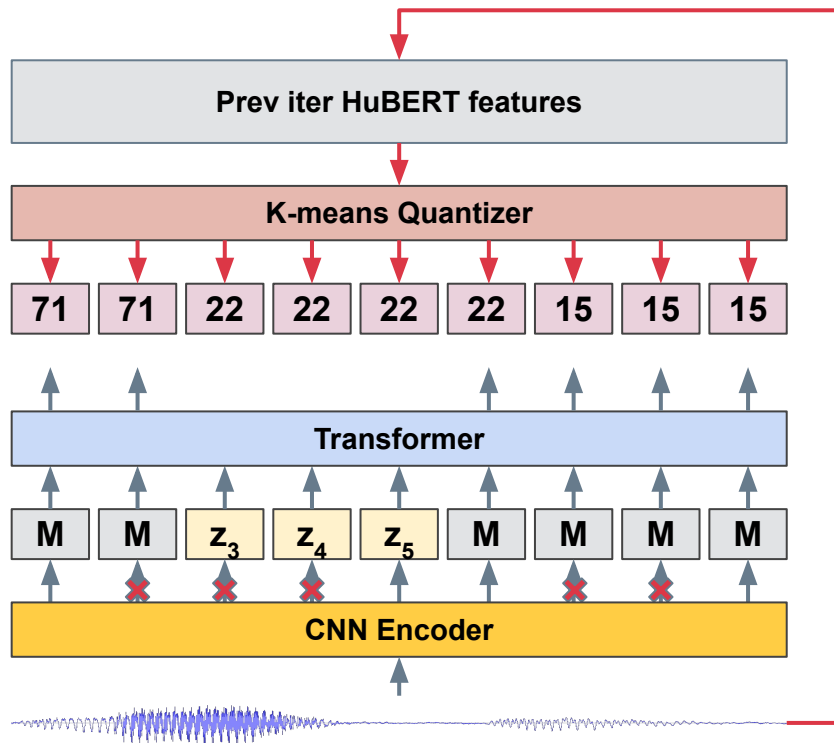
Neutral

Amused

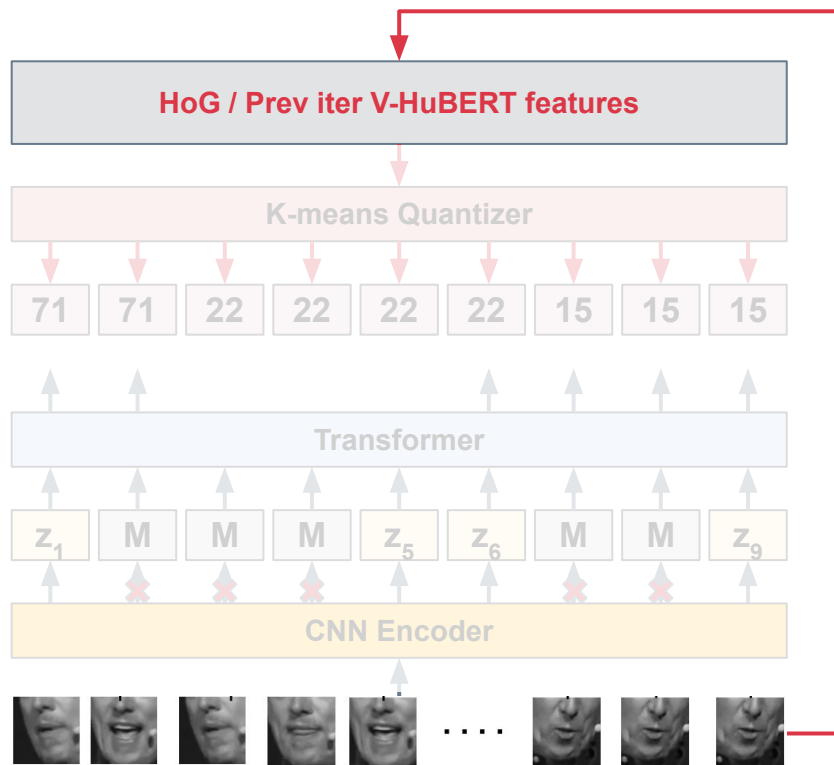


	PR	KS	IC	SID	ER	ASR (WER)		QbE	SF		ASV	SD
	PER ↓	Acc ↑	Acc ↑	Acc ↑	Acc ↑	w/o ↓	w/ LM ↓	MTWV ↑	F1 ↑	CER ↓	EER ↓	DER ↓
FBANK	82.01	8.63	9.10	8.5E-4	35.39	23.18	15.21	0.0058	69.64	52.94	9.56	10.05
PASE+ [16]	58.87	82.54	29.82	37.99	57.86	25.11	16.62	0.0072	62.14	60.17	11.61	8.68
APC [7]	41.98	91.01	74.69	60.42	59.33	21.28	14.74	0.0310	70.46	50.89	8.56	10.53
VQ-APC [32]	41.08	91.11	74.48	60.15	59.66	21.20	15.21	0.0251	68.53	52.91	8.72	10.45
NPC [33]	43.81	88.96	69.44	55.92	59.08	20.20	13.91	0.0246	72.79	48.44	9.4	9.34
Mockingjay [8]	70.19	83.67	34.33	32.29	50.28	22.82	15.48	6.6E-04	61.59	58.89	11.66	10.54
TERA [9]	49.17	89.48	58.42	57.57	56.27	18.17	12.16	0.0013	67.50	54.17	15.89	9.96
modified CPC [34]	42.54	91.88	64.09	39.63	60.96	20.18	13.53	0.0326	71.19	49.91	12.86	10.38
wav2vec [12]	31.58	95.59	84.92	56.56	59.79	15.86	11.00	0.0485	76.37	43.71	7.99	9.9
vq-wav2vec [13]	33.48	93.38	85.68	38.80	58.24	17.71	12.80	0.0410	77.68	41.54	10.38	9.93
wav2vec 2.0 Base [14]	5.74	96.23	92.35	75.18	63.43	6.43	4.79	0.0233	88.30	24.77	6.02	6.08
wav2vec 2.0 Large [14]	4.75	96.66	95.28	86.14	65.64	3.75	3.10	0.0489	87.11	27.31	5.65	5.62
HuBERT Base [35]	5.41	96.30	98.34	81.42	64.92	6.42	4.79	0.0736	88.53	25.20	5.11	5.88
HuBERT Large [35]	3.53	95.29	98.76	90.33	67.62	3.62	2.94	0.0353	89.81	21.76	5.98	5.75

Preliminary: HuBERT for audio (A-HuBERT)




Single-modal Visual HuBERT



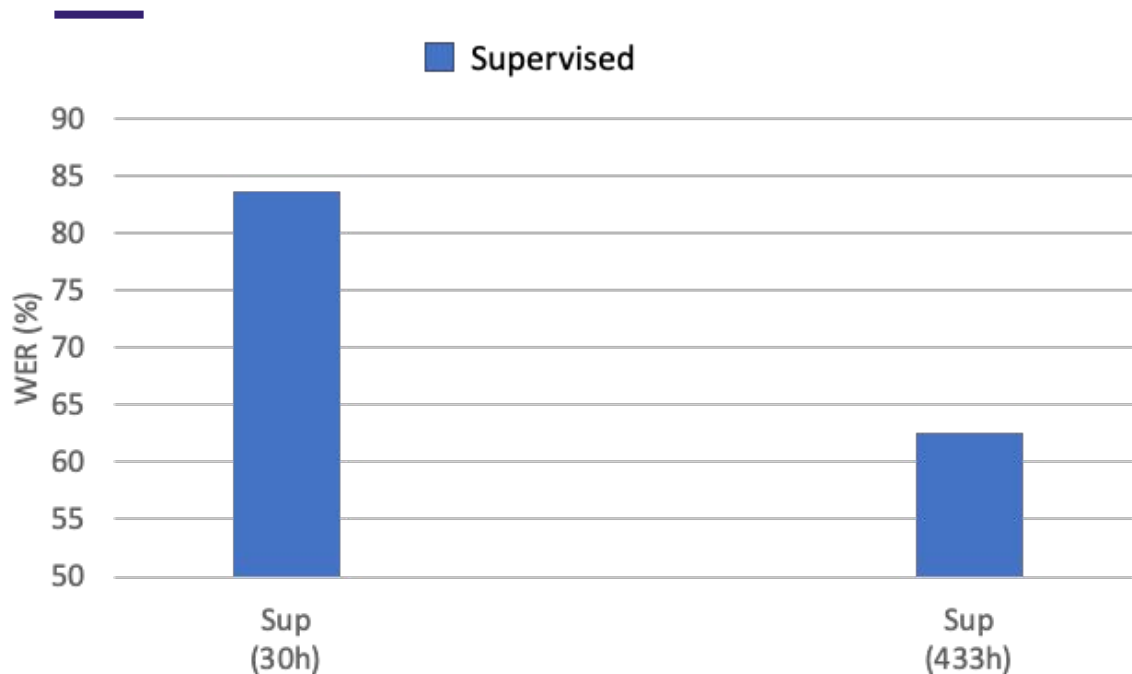
Our initial attempt, given a video stream

1. Use ~~waveform~~ images as input
2. Cluster MFCC Histogram of Gradient (HoG) features




HoG example (source: <https://www.analyticsvidhya.com/blog/2019/09/feature-engineering-images-introduction-hog-feature-descriptor>)

Single-modal Visual HuBERT



- Pre-train:
 - LRS3-433h unlabeled
- Fine-tune:
 - LRS3-30h labeled
 - CTC
- Arch: BERT Base (12L)
- Supervised baselines
 - 30h: 84%
 - 433h: 62.5%

Single-modal Visual HuBERT

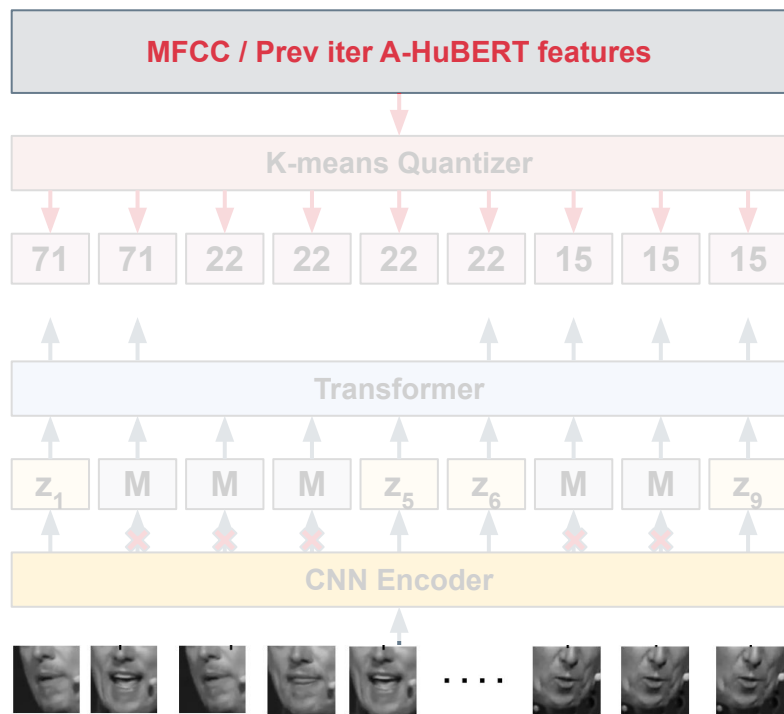


- Pre-train:
 - LRS3-433h unlabeled
- Fine-tune:
 - LRS3-30h labeled
 - CTC
- Arch: BERT Base (12L)

Limited improvement (80.1%):

- HoG cluster quality is “bad”
- Improve it with audio?


Cross-modal Visual HuBERT



Our second attempt, given audio-visual streams

1. Use lip ROI as input
2. Cluster HoG MFCC/A-HuBERT features
 - a. Leverage frame synchronicity to determine the V-A alignment


Cross-modal Visual HuBERT



Baselines:

- Sup 30h: 83.7%
- V-HuBERT, S 30h: 80.1%
- Sup 433h: 62.5%

Cross-modal Visual HuBERT




Baselines:

- Sup 30h: 83.7%
- V-HuBERT, S 30h: 80.1%
- V-HuBERT, C 30h: 69.1%
- Sup 433h: 62.5%

→ Predicting audio clusters improves visual representation learning

Cross-modal Visual HuBERT




audio cluster > video cluster

What about audio-visual cluster?

- Need AV-features
- Train HuBERT with **A+V** input
- Can we still fine-tune it for VSR?



Proposed: Audio-Visual HuBERT

- Use both audio and visual streams input
 - Mask at input independently




Proposed: Audio-Visual HuBERT

- Use both audio and visual streams input
 - Mask at input independently
 - Mask visual segments by substitution

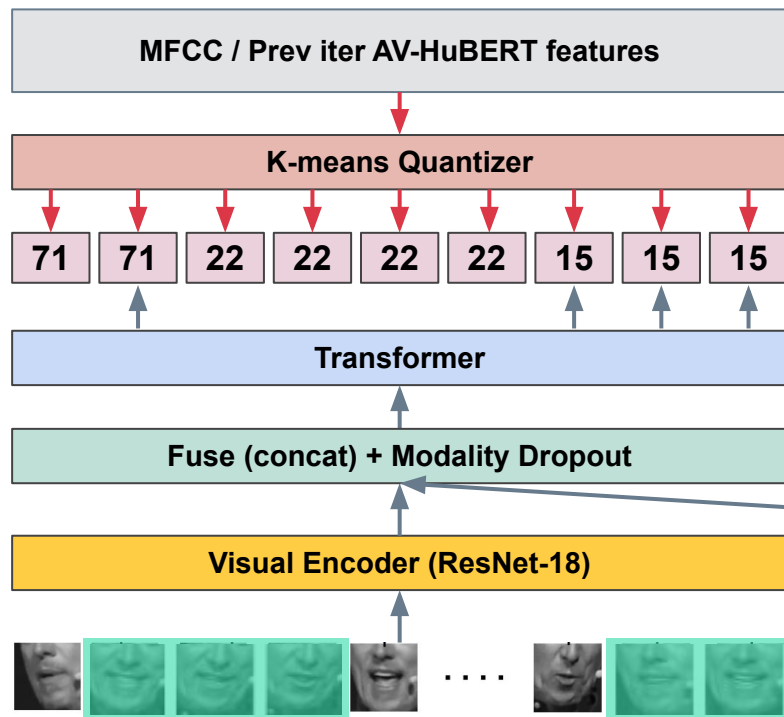


Proposed: Audio-Visual HuBERT


- Use both audio and visual streams input
 - Mask at input independently
 - Mask visual segments by substitution
- Fuse by concatenation at each frame
 - Simulate single-modal input with modality dropout (replace with 0s)




Proposed: Audio-Visual HuBERT




- Use both audio and visual streams input
 - Mask at input independently
 - Mask visual segments by substitution
- Fuse by concatenation at each frame
 - Simulate single-modal input with modality dropout (replace with 0s)
- Predict audio-visual clusters




Proposed: Audio-Visual HuBERT




Proposed: Audio-Visual HuBERT



Proposed: Audio-Visual HuBERT



Comparison with prior works




Scaling up

- More unlabeled data (1.7k hours, LRS3 + VoxCeleb2)
- Bigger model (24L BERT LARGE)
- Seq2seq fine-tuning (9L decoder)

Results:


- 30h SSL > 31K hours supervised
- Further improvement with 433hr

Comparison with prior works



Complimentary to self-training

- With 30h, 32.5% -> 28.6%
- With 433h, 28.6% -> 26.9%
- New SOTA




How Effective AV-HuBERT is for Audio-Visual Speech Recognition?

- Audio-visual speech recognition (AVSR)
 - Input: **audio+video** streams
 - output: a sequence of characters/word-pieces
 - Supervised learning: trained on (audio, video, text) tuples
- AV-HuBERT is a natural fit for audio-visual speech recognition




Supervised baselines



- BERT Large (24L enc + 9L dec)
- Supervised fine-tune:
 - LRS3-30h labeled
 - Seq2Seq
- When tested on clean data:
 - ASR: 20.6% WER
 - AVSR: 20.8% WER
- Similar performance in clean cond.


Supervised baselines



- When tested on noisy data:
 - ASR: +38.6% WER
 - AVSR: +22.1% WER


Both degrades, but AVSR is more robust

Supervised baselines




- Same trend when increasing data

AV-HuBERT results of AVSR




AV-HuBERT results of AVSR




- 35.1% absolute WER reduction
- beats 433hr-supervised model on with just 30hr labeled
- More robust to noise
 - 3.3% on clean, 7.8% on noisy

AV-HuBERT results of AVSR



- 35.1% absolute WER reduction
- beats 433hr-supervised model on with just 30hr labeled
- More robust to noise
 - 3.3% on clean, 7.8% on noisy
- Improve with more labeled data
 - 1.4% on clean, 5.8% on noisy

Comparison with Prior Work



- 19.7% absolute WER reduction compared to the prior SOTA

Conclusion

- Self-supervised learning is also very effective for audio-visual speech
- Multimodal self-supervised learning can benefit unimodal downstream tasks
- Visual information and self-supervised learning makes speech recognition more robust

Code and models available at https://facebookresearch.github.io/av_hubert

Thank you

