MEMBERSHIP INFERENCE ATTACKS AGAINST SELF-SUPERVISED SPEECH MODELS

Wei-Cheng Tseng, Wei-Tsung Kao, Hung-yi Lee
Graduate Institute of Communication Engineering,
National Taiwan University

Wei-Cheng Tseng

Wei-Tsung Kao

Hung-yi Lee

Motivation

- Self-supervised learning (SSL) models now becomes an important component of speech processing.
- To deploy SSL models to products, it is inevitable for us to make sure whether there are potential privacy issues in SSL models or not.
- In this paper, we study and propose a basic method to perform **membership inference attack** to SSL speech models.

Membership inference attack (MIA)

- Given a model and an exact datapoint, the adversary want to know whether this datapoint was used to train the model or not
- Serves as canary of more severe privacy issues

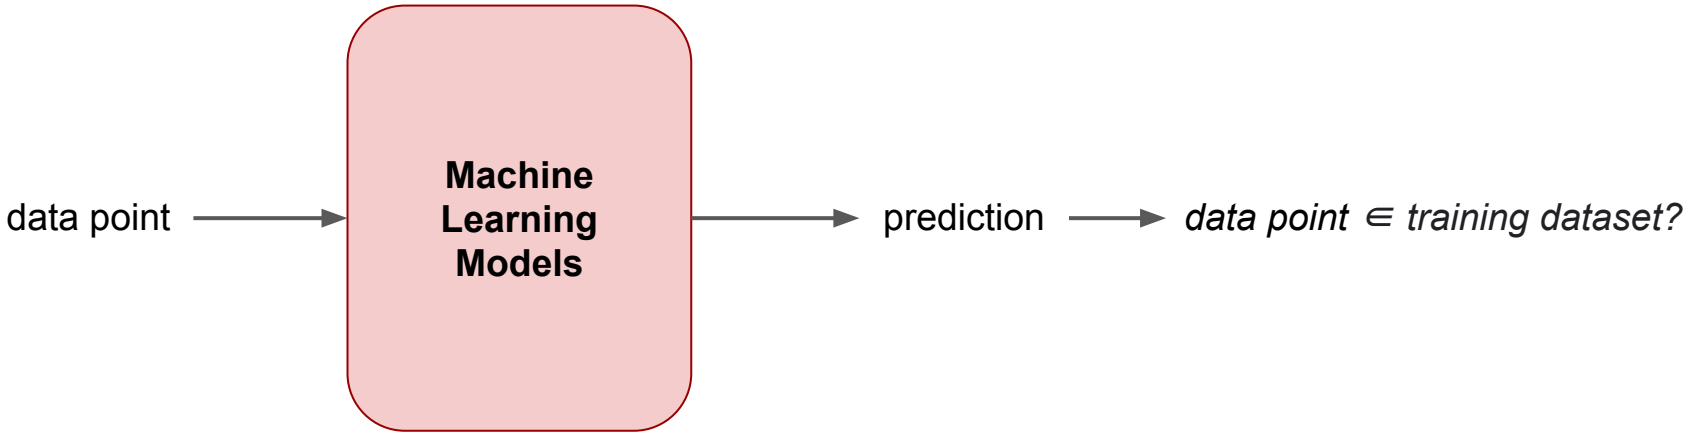

Proposed black-box MIA framework

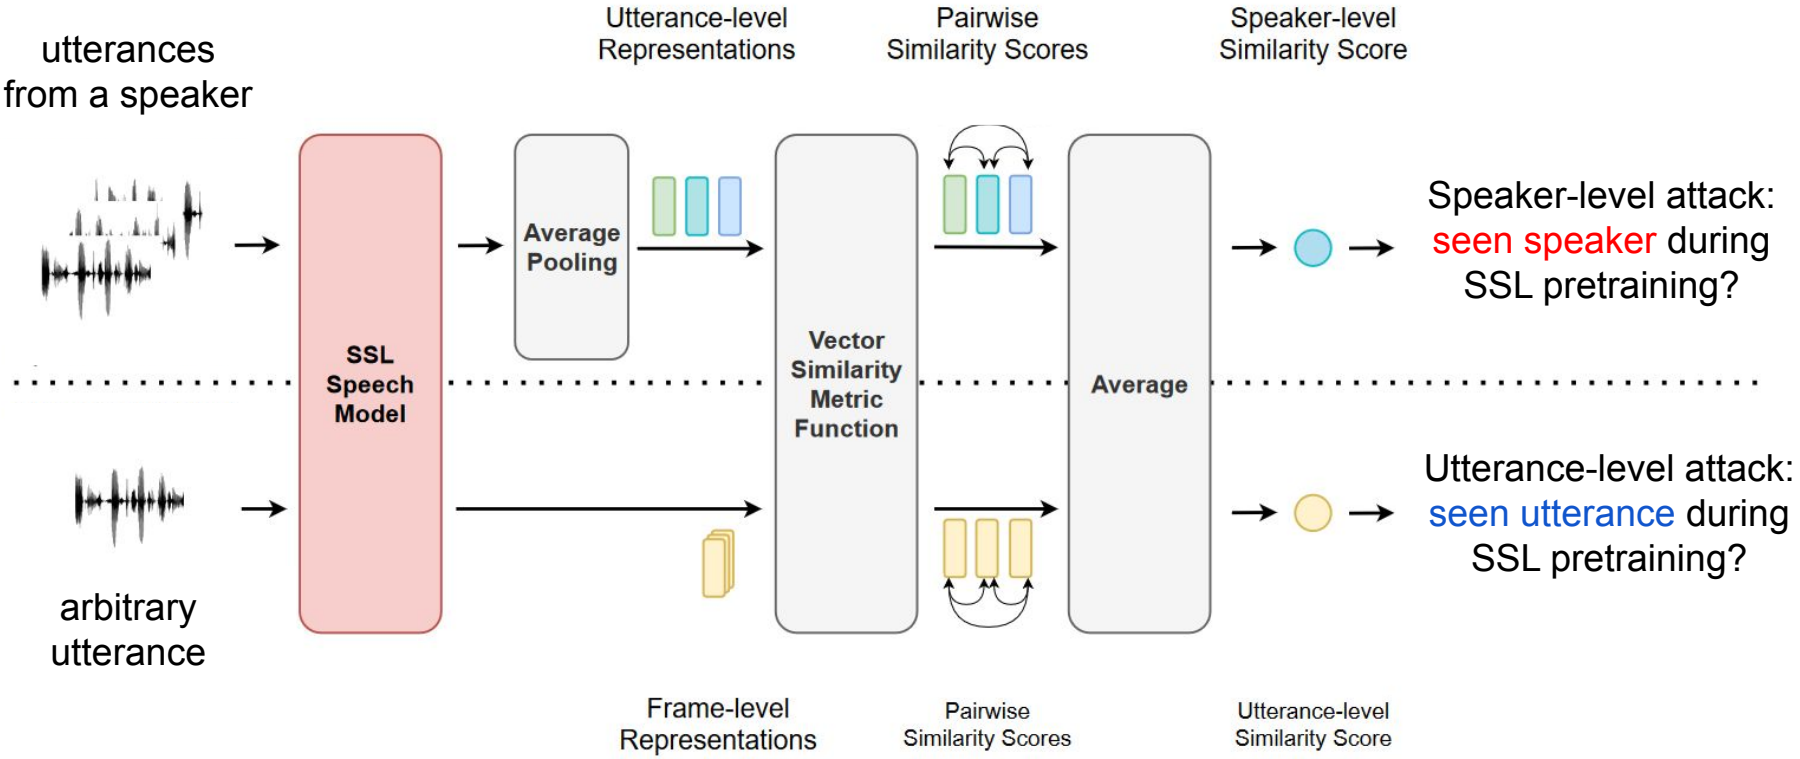

Utterance-level MIA: seen utterance during pretraining?

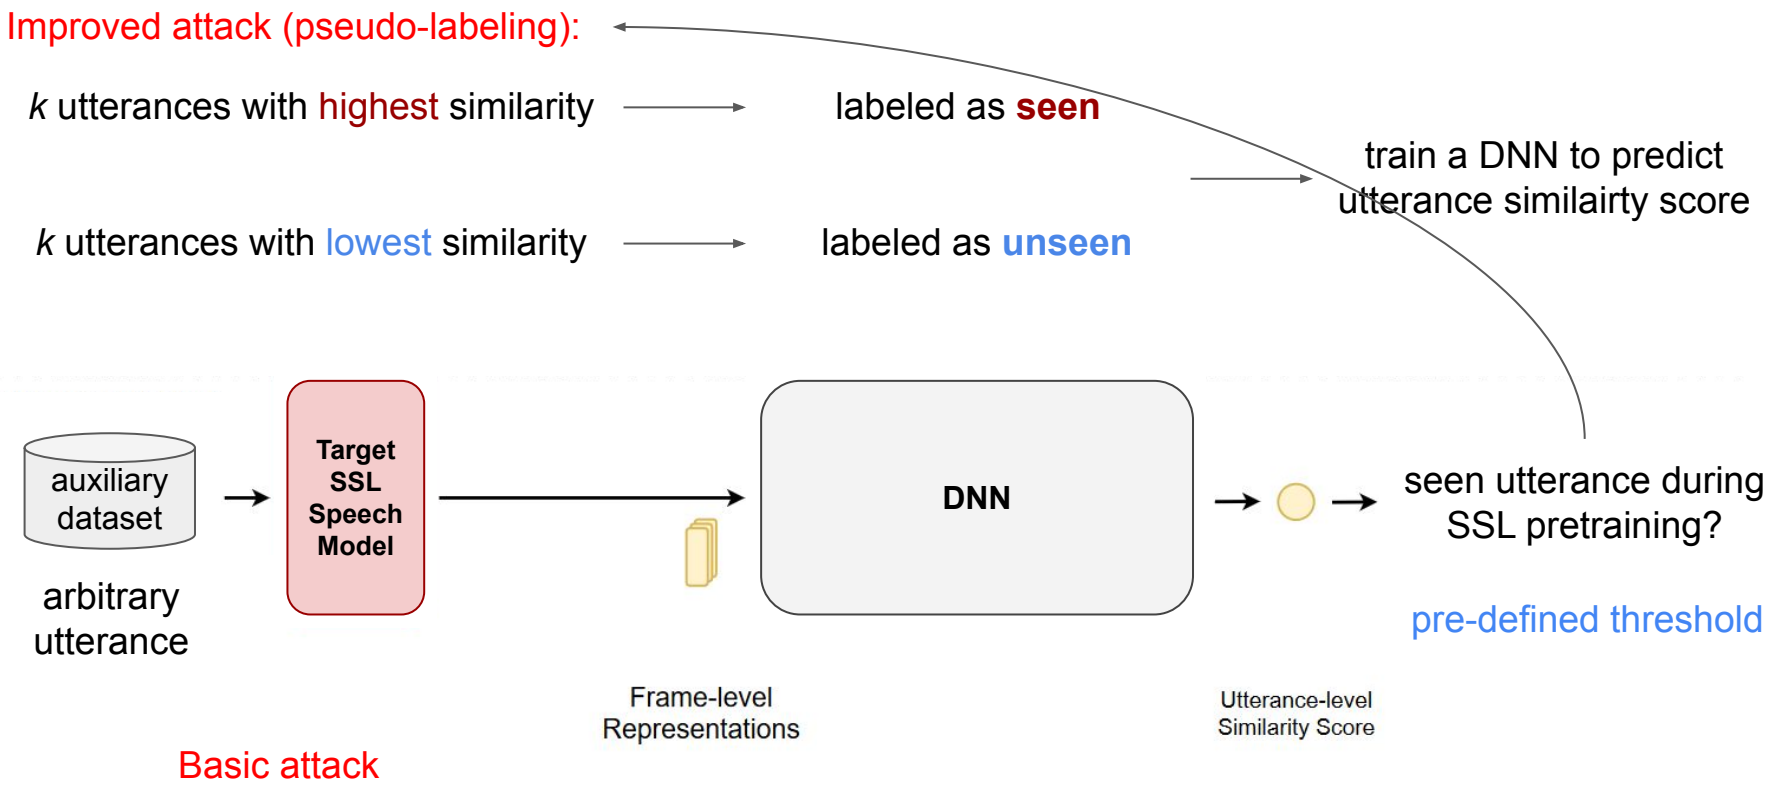

Speaker-level MIA: seen speaker during pretraining?

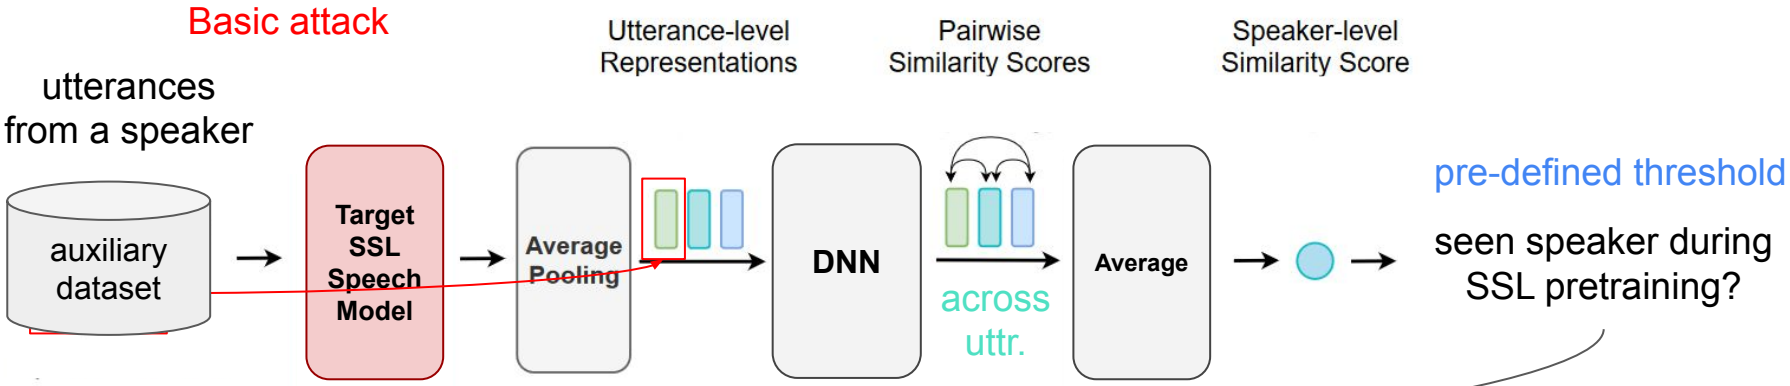

Improved attack (pseudo-labeling):

k speakers with highest similarity → labeled as **seen**

k speakers with lowest similarity → labeled as **unseen**

→ train a DNN to predict pairwise similarity score

Experiment setup

1. SSL speech models:
   - HuBERT, wav2vec 2.0, CPC and TERA
2. Dataset: LibriSpeech
   - seen: train-clean-100
   - unseen: dev-clean, dev-other, test-clean, test-other
3. Predefined vector similarity metric function:
   - utterance-level attack: 1 - (cosine similarity)
   - speaker-level attack: cosine similarity
4. *k* parameter used to train the DNN:
   - utterance-level attack: 500 (utterances)
   - speaker-level attack: 1 (speaker)
5. Evaluation: Adversarial advantage: True Positive Rate - False Negative Rate

Utterance-level MIA result

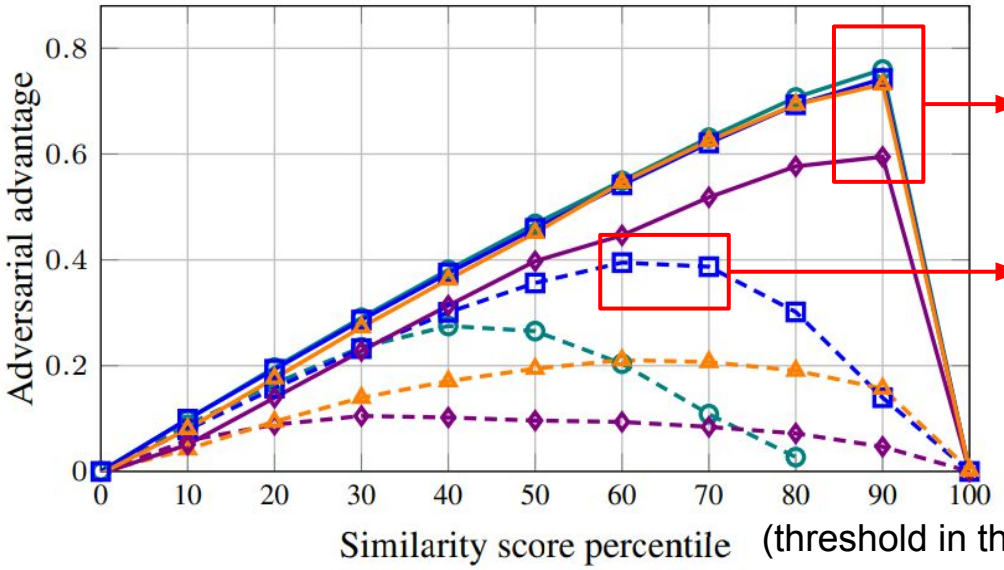

DNN significantly boosts the attack.

wav2vec2 is already vulnerable when applying basic attack

(threshold in the binary cls.)

Speaker-level MIA result

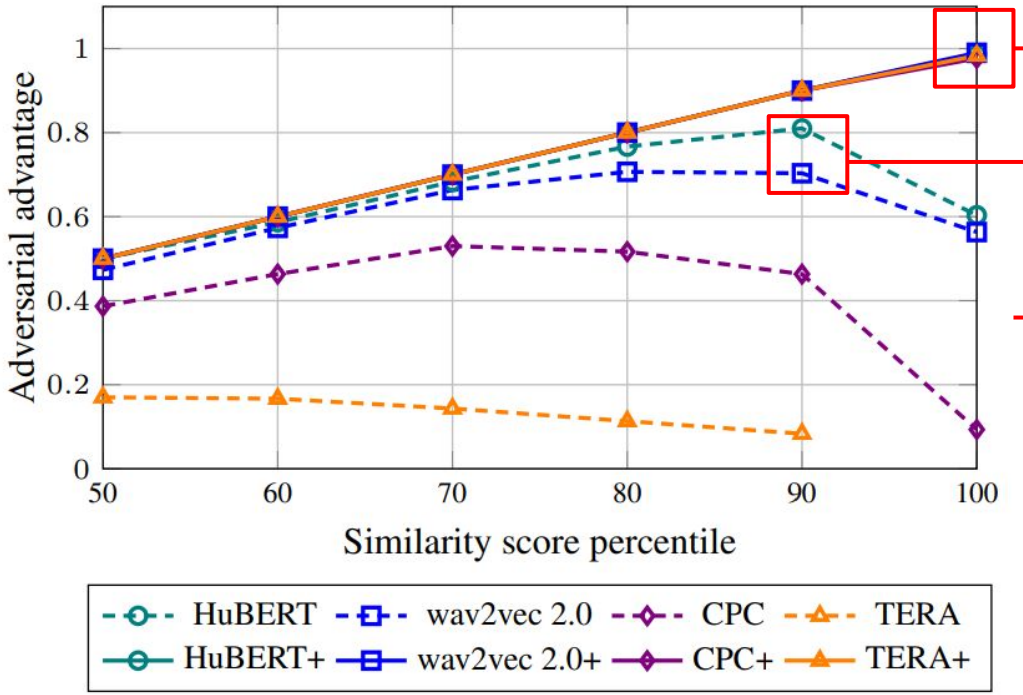

near-optimal performance for improved attack

Better models suffer from the attack more when applying basic attack

Speaker-level attack is more severe.

Choice of the auxiliary dataset

The proposed attack is robust to the choices of the auxiliary dataset.

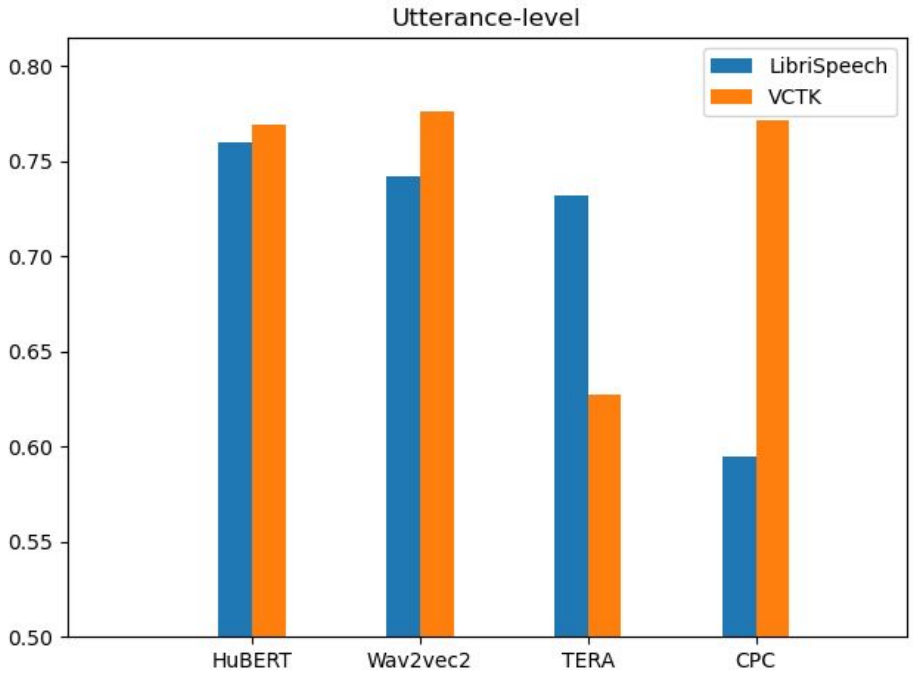

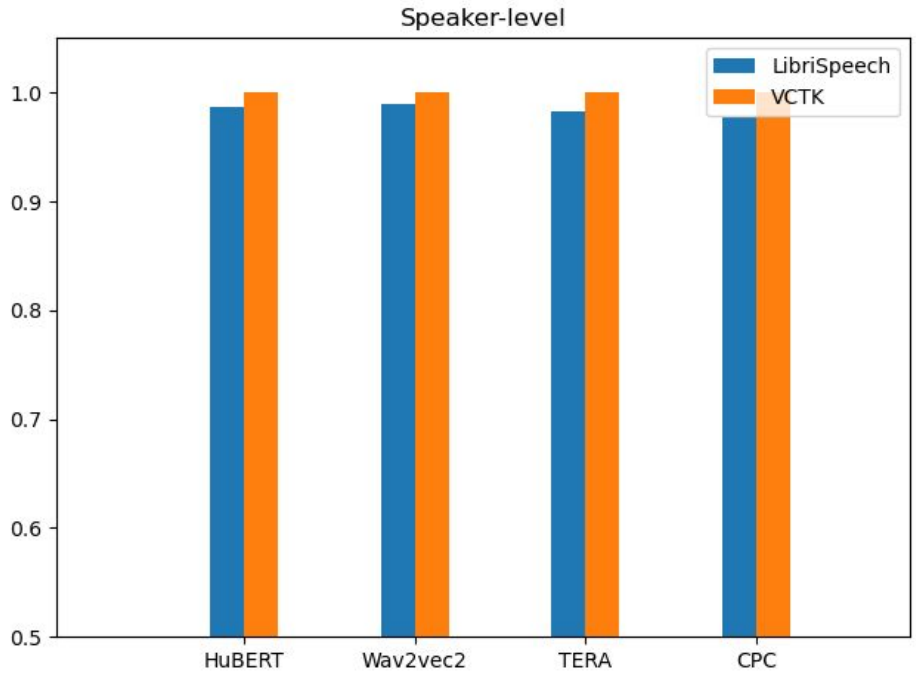

Effect of the model size

Speaker-level: Larger model -> lower attack performance

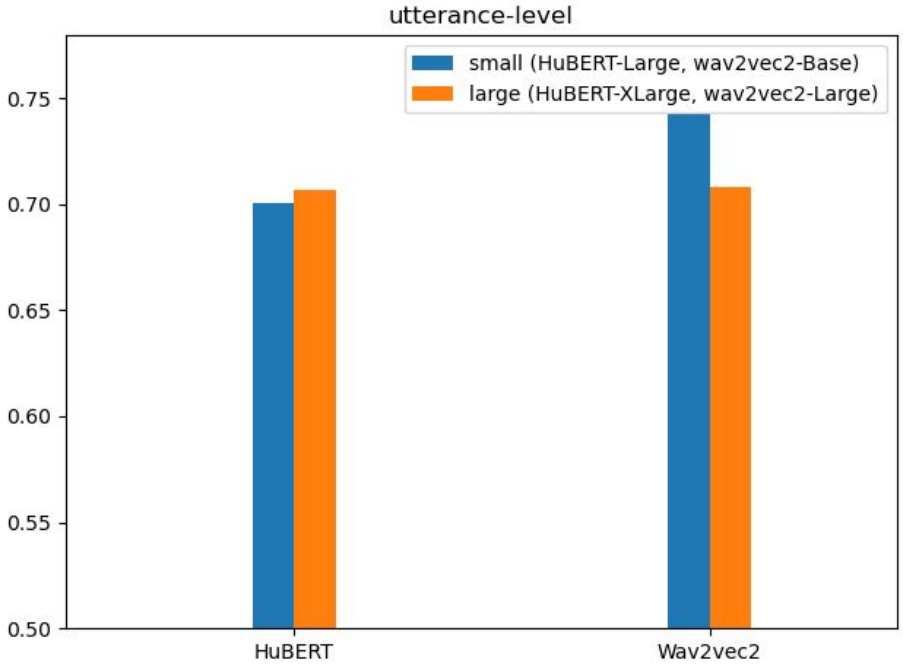

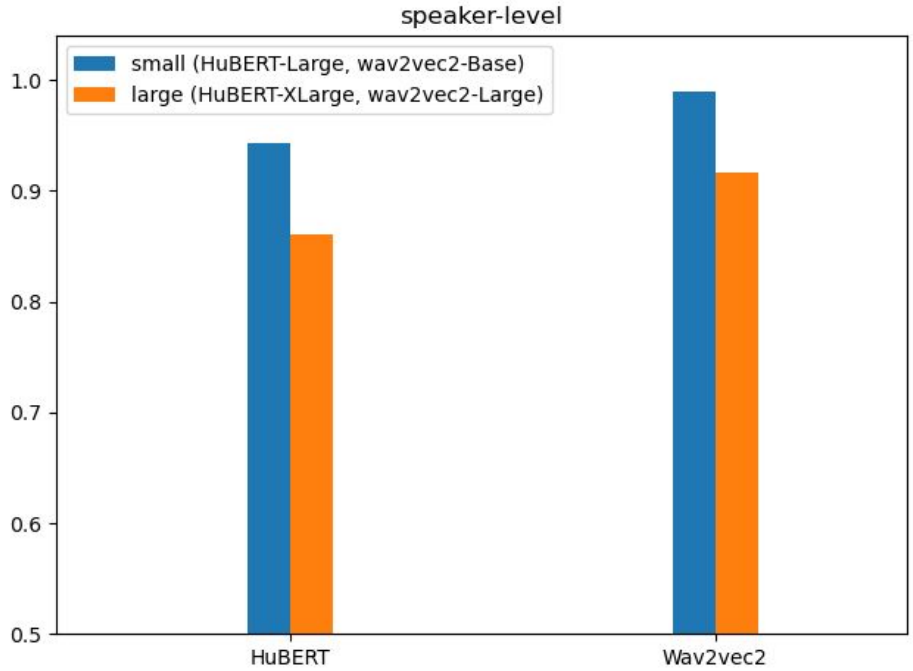

Effect of the pre-training corpus size

Smaller pre-training corpus -> lower attack performance

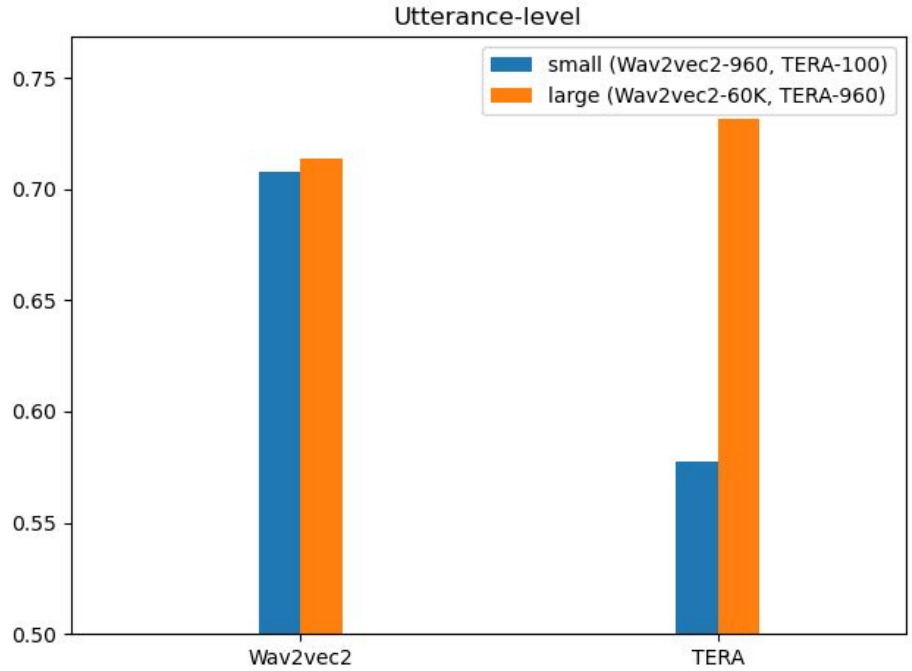

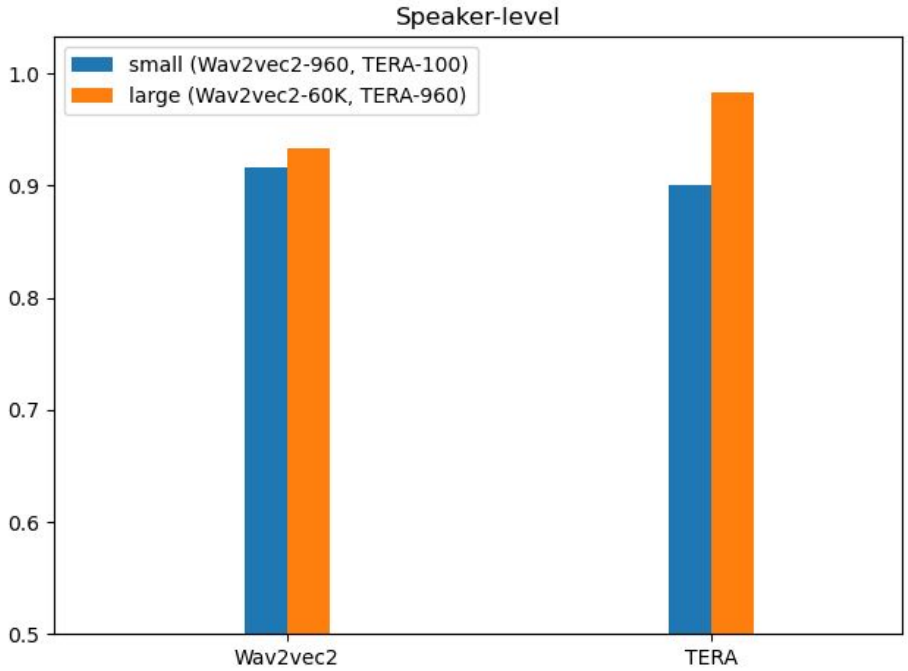

Conclusion

1. We propose the first membership inference attack against SSL speech models under black-box access.
2. SSL models are vulnerable to speaker-level and utterance-level attacks to reveal sensitive membership information.
3. The proposed attack is robust to the choice of the auxiliary datasets.
4. We also conduct ablation study on how the model size and the corpus size affect the attack performance.

Ablation study

1. Size of the model:
    a. HuBERT-{Large, XLarge} pretrained on LibriLight
    b. Wav2vec2-{Base, Large} pretrained on LibriSpeech-960hr
2. Size of pre-training datasets:
    a. Wav2vec2-Large pretrained on {LibriSpeech-960hr, LibriLight}
    b. Tera pretrained on {LibriSpeech-100hr, LibriSpeech-960hr}
3. Choices of unseen data:
    a. LibriSpeech dev and test set
    b. VCTK dataset

Effect of the model size

Utterance-level: No consistent results

Speaker-level: Larger model -> lower attack performance

Model	Model size ablation		
	Base	Large	X-Large
Utterance-level MIA			
HuBERT	–	0.7002	**0.7069**
wav2vec 2.0	**0.7423**	0.7080	–
Speaker-level MIA			
HuBERT	–	**0.9433**	0.86
wav2vec 2.0	**0.9900**	0.9167	–

Effect of the pre-training corpus size

Smaller pre-training corpus -> lower attack performance

Model	Dataset size ablation		
	LS-100	LS-960	LL-60K
Utterance-level MIA			
wav2vec 2.0	–	0.7080	**0.7134**
TERA	0.5772	**0.7317**	–
Speaker-level MIA			
wav2vec 2.0	–	0.9167	**0.9333**
TERA	0.9000	**0.9833**	–

Effect of unseen dataset

The proposed attack is robust to the choices of unseen dataset.

	HuBERT	wav2vec 2.0	TERA	CPC
LibrSpeech				
utterance	0.7598	0.7423	0.7317	0.5948
speaker	0.9867	0.9900	0.9833	0.9767
VCTK				
utterance	0.7692	0.7757	0.6276	0.7716
speaker	1	1	1	1